

Plectis: What a Stranger Can Check

Bounded public evidence from a private system

Will Cook

July 2026

Abstract

Plectis is a public software repository (an organised, copyable project folder) made from parts of a larger private system that cannot be published. It asks what a stranger may reasonably conclude from curated, runnable fragments of work they cannot see. At the version studied, each of the 88 registered components connects a claim to a fixed input, a command, saved output, a pass rule, and a stated limit. I call this arrangement a claim–evidence–limit contract and examine one component from input to conclusion. The examination exposes five gaps: a public run does not by itself establish private origin, correctness, support for the claim as worded, behaviour beyond selected cases, or the absence of risks outside a declared scan. The author chose the selection, inputs, expected answers, rules, and boundaries. The repository therefore lets a reader test claims about its published files and outputs, but cannot establish the private system’s history, whether the selection represents the whole, or whether the author’s criteria are right. I report that I developed the private system by directing AI coding tools; that account does not support the argument. A successful run supports only what the published rule warrants, which may still fall short of the prose claim.

1 The problem

I have a larger private software system that I cannot responsibly publish: it contains personal records, credentials, live browser and account information, my working notes, and unfinished tools. I would still like some of its claims to be answerable to strangers, rather than resting entirely on my say-so. *Answerable* has a definite meaning here: starting from public material alone, a stranger can make something happen that bears on the claim, can inspect the rule by which the outcome is judged, and can aim any objection at a particular public object rather than at my account of a hidden one. Neither the difficulty nor that meaning is specific to me, and this paper examines what such answerability can amount to: what may a stranger reasonably conclude from a curated set of runnable public fragments of a system they cannot see?

Plectis is my working answer: a public software repository, stored at github.com/wcook04/plectis. A *repository* is an organised project folder, usually kept with its revision history; to *clone* a repository is to copy that folder and history to another computer. Everything this paper says a stranger can check is in that repository, and the paper’s task is to describe the answer and measure how far it reaches.

I am 22 and study economics. Over about ten months I built the private system by directing AI coding agents, software tools that can edit, run, and revise code across a task. I set objectives and review criteria, accepted or rejected their output, and sent them back to revise failures; the agents wrote nearly all of the code, and I wrote nearly none of it. No other person worked on the system. This account explains why the publication problem arose and why the choices examined below are my responsibility, but it remains testimony about origin. Plectis can expose

public files and runs; it cannot independently prove how the private work was produced. The artefact does not prove the story, and the argument does not rely on it.

The limit on what a public run can establish would be the same if every line had been written by hand. Authorship might affect how often a failure occurs: for example, an implementation and its validator can share a mistake when one process drafts both. This paper compares no development methods and estimates no such frequency. The limits on inference are authorship-independent; the rates of error behind them are simply unknown.

The boundary of the evidence can be stated before any command is run:

Status	What belongs in it
Publicly checkable	The contents of a named public version, the commands listed for its components, the files those commands write, and the results of its public tests.
Reported here	The private system's history, the role of AI agents in producing it, and the private origin of published material.
Not established here	Whether the public selection represents the private whole, whether the mechanisms generalise beyond their fixed cases, whether self-supplied success criteria are independently correct or even aimed at the stated claims, and whether the private system works reliably at all.

This is a descriptive analysis of one artefact, not a statistical study. In the reporting vocabulary of Runeson and Höst [1], the object is the public repository at the named commit; each component is a *unit of analysis* (one item examined). Section 3 traces one ordinary component. Section 5 classifies all 88 by group and publication route; each registry entry also names its command, receipt, validator, and limit. Section 6 identifies the choices that remain mine. The five distinctions in Section 4 bound the conclusions. No subset of the components is used to estimate performance or represent the private system.

Reading this paper requires no programming. Reproducing its runs requires a terminal (the text window in which commands are typed), Git (the standard program for copying and versioning repositories), the Python programming language (version 3.11 or later), the `pip` installer (the standard tool for fetching and installing Python software), and basic command-line use. So that the body of the paper can stay with the argument, the exact commands, environment details, and installation notes are collected in the appendix, which is a dated record of one verification pass at the commit named there.

2 The component contract

A *component* in Plectis is one registered, testable part of the project. Its *fixture* is a fixed sample input, and its *receipt* is a saved, machine-readable record of a reference run. A *validator* is the program that applies a component's pass and refusal rules. A *commit* is a named, saved version of the repository. These terms describe files and commands; they do not grade the importance of what any component does. (Inside the repository a component is called an *organ*, and the code package is named `microcosm_core`: the project was previously called Microcosm, and several file and command names keep the old name. The generated page `ORGANS.md`, named for the old term, gives a plain description and a first command for every component.)

At commit `57ffb7f5830c` the registry, the machine-readable list of components, contains 88 entries. Every entry is required to connect the same things: a claim stated in prose; a fixture; a command; the output files the command writes; a stored receipt from a reference run; an




Figure 1: What every registered component must connect, at commit 57ffb7f5830c. The top row is words; the middle rows are mechanism; a pass from the validator is offered as evidence for the claim and must be read under the limit. Every element, words and mechanism alike, is supplied by the project. Here and in the later figures, a dashed outline marks what a stranger cannot run or observe directly; here that is the prose of the claim and the limit. The blue row names what a stranger can do to each part without asking anyone. Section 3 runs one real component through exactly this chain.

explicit rule for what counts as a pass and what must be refused; and a stated limit on what a pass establishes. Figure 1 lays the required parts out, together with the fact about them that the rest of this paper keeps returning to: every part is supplied by me, and every part is exposed to a stranger.

An assertion about private software can only be believed or doubted as testimony. A component converts the assertion into a procedure, and a procedure can fail in front of a stranger. The conversion does not make the claim true. What it changes is the form disagreement can take: a reader who doubts a component’s claim can point at a particular input, a particular line of the validator, or a particular number in the output, instead of arguing with my description of software they cannot see. A component that fails is still answerable in this sense, and a component can be answerable and trivial; answerability concerns the form a claim takes, not its truth and not its importance.

A component may also end in a declared refusal rather than an answer. The worked example below refuses an input format it does not support. Components that depend on an external tool, such as the Lean proof checker (a program that verifies mathematical proofs), must record a bounded refusal when the tool is absent, rather than reporting success. *Runnable* therefore means that the operation and its declared failure boundary are present in public form; it does not mean that every optional dependency exists on every machine. A refusal is informative only when it is declared in advance and distinguishable from a breakdown: a predeclared refusal marks the edge of what the component claims, while an unexpected crash is a defect. The two must not be read as the same kind of event. One honesty condition on refusals cannot be discharged from inside the repository at all: the public record shows each boundary as declared at the named commit, and it cannot show whether the boundary was drawn before or after an awkward case was first met. A boundary drawn afterwards turns a failure into a designed refusal in the telling. Freezing a version before an outside evaluation begins, as Section 7 proposes, is what would make that ordering checkable.

Little of the machinery is essential to the idea. What the contract needs is a claim in words; a determinate public case; a procedure a stranger can cause to run; an observable result; a rule connecting result to claim; a declared boundary of application; and a stated ceiling on inference. That the case here is a file of numbers, the procedure a command line, and the record a text document, is local convenience. The same contract could be satisfied in another field by a laboratory protocol with declared rejection conditions, or by a sealed evaluation service with published scoring rules. I make no claim about those settings; the point is only that the contract is a shape for exposing claims to strangers, and this repository is one implementation of it. Read simply as software, the whole apparatus is a test suite with a registry and unusually complete paperwork, and that description is accepted. The subject of this paper is not the machinery; it is what such machinery can and cannot make a claim answer for.

Answerability overlaps with reproducibility and transparency but is not a synonym for either. A run can repeat perfectly with no stated claim attached, and a claim can be answerable while still awaiting a good test. The design also begins from partial disclosure: it offers a bounded point of contact, not a view of the hidden system. A component exposes the published fragment to challenge. It does not make the private whole answerable.

Answerability is not the same property as computational reproducibility, although the component contract supplies materials needed to test the latter. Under the convention used by the National Academies, computational reproducibility means obtaining consistent results using the same input data, computational steps, methods, code, and conditions of analysis; such consistency does not by itself establish correctness [4]. Software testing calls the mechanism used to decide whether an output is correct an oracle, and the difficulty of obtaining a trustworthy oracle is a recognised testing problem [2]. The risk that a curated collection disproportionately exposes favourable cases is analogous to selective publication and the file-drawer problem [3]. An *assurance case* is a structured argument linking a system claim to reasons and evidence. The Object Management Group’s SACM 2.3 standard models such cases as auditable claims, arguments, and evidence, and explicitly maps Claims–Arguments–Evidence notation [6]. Plectis is not an implementation of that standard and makes no claim that the private system is safe. Its validator does not supply a full argument from evidence to claim; it applies a published rule whose adequacy remains open to dispute. The narrower object here is a public procedure attached to a claim and a limit. ACM’s policy separates *Artifacts Evaluated—Functional*, meaning an independently audited artifact found documented, consistent, complete, and exercisable, from *Results Reproduced*, meaning its main results were later obtained by another team using some author-supplied artifacts [5]. Neither is reported here. This paper claims no new theory of assurance, testing, or reproducibility. Its contribution is a worked boundary analysis of one author-curated software collection: what its public procedures expose, what still depends on the author, and what a stranger may conclude.

3 One run, examined

The component `batch8_audio_level_rms_port` computes a loudness level from short arrays of audio samples. I chose it for the worked example because it is ordinary. It needs no mathematics beyond a square root and no knowledge of the private system, and the claim it makes is small. Its value here is that every part of the evidence can be seen at once.

The reported origin is this. The private system contains a small macOS recording application, and one of that application’s functions reduces a short stretch of audio samples held in memory to a level between zero and one. The public component re-expresses that calculation in Python. The component’s code and its stored receipts name the private file it claims to re-express (`AudioLevelMonitor.swift`); a stranger cannot open that file, so the name is a report, not

evidence. What the stranger can check is everything else.

The declared calculation begins with root-mean-square (the `rms` in the component's name): square each sample, average the squares, and take the square root. It then applies the project's gain of eight and caps the result at one. Signed sixteen-bit integer samples are first divided by 32,767, the largest positive value in that representation, to bring them onto the same scale. An empty input must produce zero. An unrecognised format name must be refused rather than guessed at.

The component's public command runs it against its fixture and writes fresh output outside the repository (the backslashes split one long command across lines):

```
PYTHONPATH=src python3 -m \
  microcosm_core.organs.batch8_audio_level_rms_port run \
  --input fixtures/first_wave/batch8_audio_level_rms_port/input \
  --out /tmp/plectis-audio-rms \
  --acceptance-out /tmp/plectis-audio-rms-acceptance.json
```

At commit 57ffb7f5830c, the run produced:

Case	Input	Expected	Observed	Result
Decimal samples	0.0, 0.05, -0.05, 0.1	0.489897949	0.489897949	pass
Sixteen-bit integer samples	0, 1638, -1638, 3277	0.489892970	0.489892970	pass
Values above the cap	1.0, -1.0, 0.5	1.000000000	1.000000000	pass
Empty input	(none)	0	0	pass
Unknown format	pcm24	refusal	refusal	pass

Because the inputs and the rule are both printed above, the first row can be checked without the repository or a computer: the squares of the four samples average to 0.00375, whose square root is 0.061237..., and eight times that is 0.489897.... A reader with a calculator can recompute this expected value from the stated formula alone. For this one case the reader can act as their own *oracle*, the general name for a tester or mechanism used to decide whether an output is correct. Here elementary arithmetic supplies a standard independent of the program under test.

Almost nothing else in the collection is like that, which is why the example matters. Elsewhere the project usually supplies the implementation, the input, the expected values, and the validator through the same author-controlled process. A test assembled that way can pass for four different reasons: the implementation and the expectation are both correct; the two share a mistake; the chosen cases happen to sit where the implementation behaves; or the validator checks a weaker property than the prose claim suggests. A matching result therefore establishes repeatability under that public test. Independent evidence bearing on correctness needs a credible oracle whose expected result is justified through a route independent of the implementation, as elementary arithmetic is for the first row. Fresh inputs alone do not supply one.

The example suggests a question for any software test: where did the expected value come from, and who supplied it? The answer changes what a match can support:

Where the oracle lives	What a matching result then supports
The project itself: implementation, cases, expectations, and rule from one author-controlled process.	Repeatability under the project’s own test; the four readings above all stay open.
The reader’s own derivation from stated premises, as in the worked example’s decimal-samples row.	Bounded correctness of that one calculation on that one case.
An established external tool or reference.	Whatever the reference itself warrants, on the cases run, within the reference’s own limits.
An evaluator who chooses inputs and independently justifies expected results after a version is frozen.	Independent evidence bearing on correctness for those cases, within the oracle’s own justification and limits.
The world: an outcome observed apart from any program.	A bearing on usefulness or empirical adequacy, where such an outcome exists at all.

Most components in the collection sit in the first position, except where a reader promotes a case to the second by recomputing it, as the worked example’s decimal-samples row invites. Components that invoke a named external tool occupy a hybrid of the first and third positions: the project still chooses the case and the rule, while the tool supplies a result that inherits its own limits. The routes in Section 7 describe ways to move consequential choices out of the project’s control.

The refusal row records a designed boundary, not a failure: an input naming an unsupported encoding (`pcm24`, a format this component does not handle) must be rejected with an error, and is. The stored reference for the whole table is `batch8_audio_level_rms_port_validation_receipt.json` under `receipts/first_wave/batch8_audio_level_rms_port/`, so a fresh run is compared against a fixed record rather than against memory. A stored record by itself does not prove that the stated command produced it; its use is that anyone can generate a fresh one and compare the two.

The limit is part of the component. This run checks one public Python calculation over synthetic sample arrays (arrays composed for the test, not captured from a device). It does not run the original application, start an audio session, request microphone permission, or capture sound, and it cannot establish that the named private file is where the calculation came from. Nor does the table show that the task was difficult (it was not), that the component is useful elsewhere, or that the remaining entries resemble it. What it shows is that one assertion now has a form in which a stranger can rerun it, dispute its rule, perturb its input, or replace my expected value with a better-derived one.

4 Five distinctions

Five distinctions organise the rest of the analysis. Each separates an observation from a conclusion it may seem to support, then names the further evidence needed to connect them. Individual component limits are instances of these distinctions. Figure 3 sets them side by side.

Public execution versus private provenance

A reader can establish what a named public commit contains and what its commands do. No public run establishes where the material came from. Plectis keeps records at that boundary: copied files carry public import records naming their claimed sources, together with SHA-256

fingerprints, fixed-length digests used here to detect changes. Different files can in principle share any fixed-length digest; the standard defines the algorithm, not a collision-free guarantee [8]. Matching fingerprints show that a public file agrees with the project’s public record of it. But I control both sides of that comparison, so agreement is evidence of internal consistency, not of origin. The same holds for the audio component’s named source file: the name is an entry I wrote. Origin claims could only be strengthened by something a public repository cannot supply about itself, such as commitments published before the fact, or a trusted person permitted to examine both sides. Until then, provenance is reported.

Repeatability versus correctness

A run that reproduces the stored receipt shows that the public code, on the supplied input, in a sufficiently similar environment, produces the same result again. Correctness is a different property: it needs a standard for what the result should have been, and a standard is only worth something if it does not merely restate the implementation. The worked example makes the difference concrete. Its first row has an oracle independent of the project, namely elementary arithmetic; most components have only the project’s own expected values and validator. Passing tests anywhere in Plectis should be read at that strength and no more.

A validator’s rule versus the stated claim

A pass is a verdict from the validator, and a validator checks exactly the property written into it, which can be narrower than the claim as worded. A component whose prose speaks of handling a class of inputs may in fact verify that one output file exists and matches a stored shape; a rule can hold while the sentence it stands under goes untested. The two previous distinctions do not cover this gap: a result can be repeatable, and correct as far as the checked property goes, while the checked property is beside the stated point. Plectis makes the gap inspectable rather than closed. Every validator is public, so a reader can put the rule beside the prose and object where the rule is weaker, and the review in Section 8 makes exactly that comparison its central step. Since the same author-controlled project produced both the claims and the rules, agreement between them records one judgement twice, and no count of passing components can substitute for reading one rule against one sentence.

Selected cases versus general behaviour

This distinction applies twice. Within a component, the fixture’s cases are a few points in a large space of possible inputs, and behaviour elsewhere is untested. Across the collection, the published components are a selection from a private whole, and the selection was mine. I chose which mechanisms could be given public form, which inputs to freeze, which expected values to store, and which pass rules to enforce. Nothing in the public record shows how many candidates there were, what was excluded beyond the privacy rules, or whether the collection resembles ordinary work in the private system. A collection can consist entirely of genuine items and still give an unrepresentative picture of what it was drawn from. This is analogous to selective publication and the file-drawer problem, not the same process [3]. Figure 2 draws the situation: public checking reaches everything to the right of the boundary, and the pool the selection was drawn from stays on the left.

The registry’s count should be read in that light. A registry of 88 entries defines the published population exactly, which is what a later evaluation would need in order to sample from it and report failure rates against a fixed denominator. The count is an inventory, not a score, and




Figure 2: The selection problem. Each square stands for a private mechanism; the darker ones were given public form, and each, unpacked, has the anatomy of Figure 1. The grid trails off because its true extent is exactly what the public side cannot know. The count of squares on the left, the rule that picked the darker ones, and the resemblance between the two sides are all invisible from the right. The collection supplies directly checkable evidence about the published files and their behaviour under named tests, but no basis by itself for generalising to the hidden whole.

even the unit of counting is a choice I control, since one mechanism could have been registered as several components, or several as one.

Risk reduction versus guarantee

The publication process runs checks whose honest description is that they reduce specific risks. Automated scans search copied material for specified patterns of credentials, personal information, and private file paths, and record what was omitted or rejected; a scan establishes that the scanner did not find the patterns it was designed to find, and nothing stronger. Pages such as the component map are generated from the registry, and tests compare page and registry, which prevents the particular failure of prose drifting away from records; a generator and its records can still be wrong together. The documented `tour` run takes the public checkout as its explicit input and does not automatically discover a neighbouring private repository. The release-export contamination check is an explicit exception: when a copy of the private project folder is available, it may compare the contents of public source files with private files to detect leakage. None of this amounts to a guarantee that no secret, no identifying detail, and no inconsistency survives. Publication remains a judgement about acceptable risk, and I made it.

The provenance and risk-reduction distinctions pull against each other in a way worth noticing. Establishing provenance would require someone to examine private material, and the privacy constraint that motivates the whole design forbids exactly that examination in public. Some claims are therefore permanently in the reported category, short of a confidential review. The honest treatment is to label them, not to dress them up as public proof through bookkeeping that one maintainer controls.

The five distinctions share one structure. What a stranger observes sits inside the published frame: a run completed, a record matched, chosen cases passed, a scan finished. The conclusion the observation invites lies outside the frame, and reaching it needs a premise the observation




Figure 3: Five conclusions a matching run invites and does not license. Each dashed arrow needs a further premise that the observation itself cannot supply, and the tags follow the order of this section’s distinctions: origin for provenance, standard for correctness, meaning for the validator’s rule against the stated claim, reach for selection, and coverage for risk. The contract’s contribution is not to cross these gaps but to make them nameable, one component at a time.

does not contain, one premise per distinction: that the public object descends from the private one; that the expectation is what the answer should have been; that the rule tests what the words say; that the chosen cases stand for the unchosen; that the risks searched for are the risks there are. The component contract does not supply these premises, and cannot. What it does is make each missing premise nameable, and a named gap can be argued about, priced, or deliberately left uncrossed, where an unnamed one is crossed by default. Nor is five a closed count. The same shape has already appeared twice outside the list: a refusal boundary observed today does not show when it was drawn (Section 2), and a number observed at one commit does not describe the next (the appendix). Five is how many this collection most needs a reader to hold, not how many there are.

5 What the collection contains

At commit 57ffb7f5830c the 88 components fall into 7 groups:

Group	Count	What its components do
Getting started	2	Show a new reader where to begin and guide a first inspection.
Project mapping	12	Map files and written rules, find relevant material, and route a question to the part of the project that owns it.
Computer-checked mathematics	20	Prepare, repair, or check bounded steps used in formal mathematical work.
AI reliability and safety tests	20	Replay fixed cases involving hostile instructions, software boundaries, stored information, tool use, and benchmark rules.
Research-method tests	9	Exercise fixed cases involving replication, conflicting forecasts, explaining model behaviour, simulation, and laboratory safety.
Public-copy maintenance	20	Check copied source, generated pages, publication boundaries, and disagreement between records and derived files.
Work recovery	5	Record unfinished changes and test how work resumes after interruption.

The groups do not share a subject; computer-checked mathematics has nothing in particular to do with audio levels or with recovering interrupted work. Their unity is the contract of Section 2: these are the places in the private system where a mechanism could be given a bounded, runnable public form. The collection is organised by how its claims are checked, not by subject.

The registry also records how each public form was produced. Four routes cover all 88 entries, and each is a different trade between fidelity to the private original and what a stranger can check:

Route	Count	What it preserves	What it cannot show
Record-checked copied source	21	The exact text of non-secret source files, checked against a public import record.	That the recorded origin is real; record and file share one maintainer.
Public reimplementation	27	The behaviour of a bounded calculation, exercised on fixed cases.	Any relation to the private original beyond my report; the original text is absent.
Direct or external-tool run	22	A local computation, or a named external tool's result, captured with its context.	Behaviour where the tool or environment differs; the result is as strong as the tool and the case.
Contract checker	18	Whether a public input satisfies a declared structural rule.	Whether the declared rule is the right rule to require.

The copied-source route preserves the most text and demonstrates the least behaviour. A reimplementation demonstrates behaviour while surrendering textual identity; the worked example is one of the 27 in that route. An external-tool run inherits whatever confidence the tool deserves, along with the tool's limits. A contract checker is worth exactly the contract its author chose to require. A reader who wants to know which trade a given component made can read its registry entry, which names its command, its fixture, its receipt paths, its evidence route, and its limit; the plain description of what it does lives on the generated component page named in Section 2.

Two cautions govern any reading of these tables. The first caution is arithmetic: components are not independent witnesses. Entries share code, fixture style, validator authorship, and my judgement; several entries can descend from one private mechanism, and where they do, their verdicts tend to move together. 88 entries are 88 registered claims, not 88 separate confirmations of anything. The same discount applies across the routes: the four routes are four different ceilings on inference, not four corroborating methods, and agreement among the pages, registries, and receipts that I generate from one source of truth is one witness, recorded several times. Machine-checked consistency among those records raises the cost of accident, not the number of witnesses.

The second caution is association. Group names such as computer-checked mathematics and AI reliability and safety tests borrow gravity from serious fields, and the borrowing happens in a reader whether or not any sentence asserts it. The corrective is to restate each group as the weakest property its validators actually check, which is often that a fixed case replays or that a record matches a declared shape, and then see how much of the impression survives. The names are for finding components, not for weighing them.

6 The author's hand

All five distinctions return to control. At the commit this paper describes, I hold every consequential choice on both sides of every test. I chose which private mechanisms were candidates, which were selected, and how finely each was divided into components, and so what the count counts. I chose which of the four routes each took, which inputs were frozen, what the expected values are, what each validator checks, and where each refusal boundary sits. And I chose which group each entry is filed under, which failures this paper explains and in what tone, which single component the worked example honours, and where the publication boundary lies. Somebody had to make each of these choices, and the privacy constraint meant it was me. None of that is misconduct.

But that concentration of choice permits a failure mode, and it involves no false statement anywhere. Select the mechanisms that show well. Divide the favourable ones finely. Narrow each claim until its fixed cases satisfy it. Route awkward material through the cheapest form of checking. Declare as refusals whatever would otherwise fail. File small replays under weighty names. Then let the counts, the categories, and the machinery of records accumulate into an impression of scale and maturity that no individual sentence asserts. A document can be cautious sentence by sentence while its arrangement still overstates, because a reader's impression forms on what is counted, named, worked through, and left out, as much as on what is claimed. This paper is inside its own arrangement, so declaring the risk does not discharge it; nothing in the text can establish that the text is not doing what this paragraph describes.

Making disagreement local can also displace the larger question. A reader can settle an argument about one input or one rule while leaving untouched whether the selection is fair, whether the whole is sound, or what any of it is for. An objection about selection is not answered by the successful execution of any number of selected items. The registry hosts the small arguments; it does not settle the large ones.

Two defences need no cooperation from me. Read any component at the weakest property its validator checks, as Section 5 proposes. Treat every count as an inventory, and every agreement among my records as one witness. A third correction is a standing discount rather than an action, because it concerns what no reader can inspect: a mechanism survives selection when it can be made bounded, deterministic, and runnable in isolation, so a collection like this over-represents pure calculation, replay, and record-checking, and under-represents stateful, entangled, long-




Figure 4: Five independent evidence routes. Above each route, in blue, the choice that leaves my hands; below it, in grey, what that route still cannot show. The routes are not cumulative stages. At the commit this paper describes, none has happened.

running behaviour. The published picture tilts towards what publishes well before any question of honesty arises.

Disclosure alone cannot remove this concentration of control. NASA’s standard separates technical independence (evaluators did not develop the system), managerial independence (they choose what and how to assess), and financial independence (funding is controlled outside the development organisation) [7]. Plectis has had no such evaluation. Within the five gaps examined here, the narrower principle is that evidence for a claim becomes less dependent on me as the relevant choices stop being mine. The principle is local to those five gaps. Section 7 lists routes that address the dependencies this paper has identified.

7 What stronger evidence would look like

A passing run shows that named public code produced a stated result from a supplied input in a stated environment. Applying Section 4: the runs reported here do not establish private provenance, representativeness, correctness beyond self-supplied criteria, or anything about usefulness, security, originality, or the reliability of the private system as a whole. The computer-checked mathematics components prove no more than their named public fixtures, and no other repository’s size or subject matter is used as evidence here. This note does not report an independent, randomly selected audit over fresh inputs; nothing of that kind has happened yet.

Five independent routes could move different boundaries. They are ordered below for exposition, not by strength or prerequisite: they may occur in any order, and progress on one route does not imply progress on the others. Figure 4 shows which choice each route transfers and which claim it still leaves open. The shared question is who controls the consequential choices bearing on that claim.

Independent repetition means a person with no involvement in the project reruns the supplied commands in their own environment and publishes what happened, including failures. That would reduce dependence on my machine and test whether the published instructions suffice elsewhere. It would not by itself establish correctness.

Evaluator-controlled selection means the evaluator chooses components from the full registry, without substitution when one proves awkward, and original failures stay on the record even if I later repair them. The units themselves are open to the same challenge: an evaluator who finds that two entries are one dependent mechanism, or that a boundary excludes the difficult part

of a task, is disputing the registry, not misusing it. That addresses curation within the public collection. It cannot address whether the collection represents the private system, because the private pool stays invisible.

Fresh inputs with independently derived expectations can move repeatability towards evidence bearing on correctness: inputs chosen after a version is frozen, and expected results worked out through some route that does not pass through my code, whether independent calculation, a second implementation, or an external reference. The result would support bounded claims only to the extent that the derivation is credible and independent and the cases are adequate; general correctness still would not follow.

Comparison on outside tasks is what a claim of advantage over ordinary alternatives would need: someone bringing a task that was not designed around the project's own categories and comparing the component against an ordinary alternative, which for several components would be a short conventional script. Nothing in this paper makes that comparison.

A *provenance review*, among the five routes shown here, bears most directly on the origin story: a trusted person examining selected private material against the public collection under confidentiality. Prior public commitments could provide another kind of provenance evidence. Even completed, this review would transfer trust rather than remove it: public readers would be relying on the reviewer's access, competence, and report instead of on mine, and the review itself could not be published. The privacy cost may reasonably never be paid. If it is not, the origin story remains testimony, and this paper is written so that its argument does not depend on it.

How the record should age

The routes above may expose failures, and the value of the record then depends on what a failure does to it. The rule this project adopts is that repairs add facts and do not subtract them. The appendix records that at the pinned commit three public tests failed, and that a later commit repairs one of them; the repair is a dated fact about a later version, and it does not subtract the first fact about this one. The same rule holds above the level of code. If outside criticism shows a validator too weak or a claim too broad, the revision should be legible as a revision, carrying its date and its cause, rather than the project appearing always to have meant the narrower thing. A test rewritten after its outcome is known is a different test and may be better. The known case supplies regression evidence; fresh confirmation requires held-out cases whose outcomes were not used to design the repair. Disappointment is in scope: an evaluation that concludes a component is trivial, or no better than a short ordinary script, is the machinery working, and its conclusion belongs on the record with the same standing as a pass. An outside evaluator's report, where one comes to exist, should remain a separate document under the evaluator's control, cited rather than absorbed. Where the repository's future conduct falls short of this paragraph, the shortfall is itself a finding, and Section 8 says where findings go.

Each route moves one class of consequential choice out of my hands. More components or records produced while I still choose every case and rule do not make that transfer. These five routes answer the gaps identified here; other claims may require other forms of evidence.

8 Conclusion

Plectis does not make the private system answerable as a whole. It puts particular public claims within reach of challenge by publishing a fixed case, procedure, output, rule, and limit. A pass is no stronger than its rule or than the route by which the expected answer was chosen.

Adding more author-controlled components enlarges the inventory; it does not alter who chose the cases or rules. For the dependencies exposed here, stronger evidence requires at least one consequential choice to leave the author’s hands.

A first review. Clone the repository and, for this paper’s numbers, check out the commit named in the appendix. Run `plectis tour`, then choose a claim you doubt. Read its description and limit, run its command, compare the new output with the stored receipt, and inspect the validator. Ask whether the pass rule bears on the prose claim; perturb the input where practical; and, where a refusal boundary does much of the work, ask when that boundary entered the history. Keep discrepancies and report them through `CONTRIBUTING.md`. One successful run supports only the conclusion licensed by the published rule. Broader conclusions need the evidence in Section 7, and most of it does not yet exist.

A A dated reproduction record

Unless a paragraph gives a later date, everything in this appendix is an observation about one commit, made on 17 July 2026 in one environment: macOS 15.6.1 (arm64) with Python 3.13.7. The commit’s full identifier is

57ffb7f5830cc446a2cbd8083d5d80bcab2af563

abbreviated `57ffb7f5830c` elsewhere. Later commits and other machines may give different results; the body does not depend on those differences.

Names. The repository is github.com/wcook04/plectis. As Section 2 notes, the code package inside it is named `microcosm_core`. The installed command is `plectis`, and an older command name, `microcosm`, still works as an alias.

Getting the code and running without installing.

```
git clone https://github.com/wcook04/plectis
cd plectis
git checkout 57ffb7f5830c
PYTHONPATH=src python3 -m plectis tour --format text .
```

The final full stop in the `tour` command means “the current folder”. The `git checkout` line pins the clone to the exact state this appendix describes; omit it to run against the newest version instead, and expect the dated numbers below to differ if you do.

In a clean clone at this commit, the `tour` read 5,018 project files, found five reading orders, chose the README-first route (`readme_onboarding_route`), wrote 21 files under `.microcosm/`, and left the examined source files unchanged. Each recorded observation names its evidence, event, and limit. The generated directory can be deleted and rebuilt. The README records a different file count from another day; both counts describe one dated run, not a stable measurement of the project.

Installing.

```
python3 -m pip install .
plectis tour --format text .
```

The runtime needs no third-party library; `pip` may download `setuptools` to build. No check enforces the claim that components avoid network and model calls; inspect the source. Before installation, the bootstrap script checks the shipped inputs and limits.

The worked example. The command in Section 3 writes under `/tmp/plectis-audio-rms`; rerunning it overwrites those files. At this commit every displayed case passed, and the fresh receipt and values matched the tracked copies. One field in the larger tracked result is stale: it reports one copied private module, while the pinned manifest records zero because public code replaced the omitted module. The calculation is unchanged, but the discrepancy remains in the record.

A second interpreter check. On 18 July 2026, a clean clone on macOS used Python 3.12.7. A network-disabled install could not obtain `setuptools>=77.0.3`; the host’s `setuptools` 80.9.0 allowed a local build. The tour and worked example passed with matching values and no added runtime library. This covers 3.12.7 and 3.13.7, not the full Python 3.11-or-later range.

Project-wide checks. At this commit, `make check` loads the component registry and finishes in under a second. The Lean trust checker scans 58 proof files for placeholders, project-granted axioms, evaluation outside the checker’s core, unsafe or partial declarations, and raised limits. Both checks passed. The 58 files and 20 mathematics components count different things. `make ci` also installs, tests, and runs two examples; continuous integration runs the same checks. They concern the public checkout, not the private system.

The full public suite was not green: three tests failed, one in the front-door profile and two in older documentation and record checks. A later commit repairs the profile defect, but this record still stands. The failures remain here, as Section 8 asks readers to keep their own failed runs.

The paper itself. `scripts/check_public_system_paper.py` recomputes counts, checks the named commit and worked example, guards the main cautions, and fails when the paper and repository disagree. Agreement between records with one maintainer is internal consistency, not external truth. The checker reads the pinned commit rather than silently comparing a historical appendix with the current checkout; it checks current prose and citations separately.

References

- [1] Per Runeson and Martin Höst. “Guidelines for Conducting and Reporting Case Study Research in Software Engineering.” *Empirical Software Engineering*, 14(2):131–164, 2009. [doi:10.1007/s10664-008-9102-8](https://doi.org/10.1007/s10664-008-9102-8).
- [2] Elaine J. Weyuker. “On Testing Non-Testable Programs.” *The Computer Journal*, 25(4):465–470, 1982. [doi:10.1093/comjnl/25.4.465](https://doi.org/10.1093/comjnl/25.4.465).
- [3] Robert Rosenthal. “The File Drawer Problem and Tolerance for Null Results.” *Psychological Bulletin*, 86(3):638–641, 1979. [doi:10.1037/0033-2909.86.3.638](https://doi.org/10.1037/0033-2909.86.3.638).
- [4] National Academies of Sciences, Engineering, and Medicine. *Reproducibility and Replicability in Science*. 2019. [doi:10.17226/25303](https://doi.org/10.17226/25303).
- [5] Association for Computing Machinery. “Artifact Review and Badging—Current,” version 1.1. Accessed 18 July 2026. [ACM Publications Policy](https://www.acm.org/publications/policies/artifact-review-and-badging-current).
- [6] Object Management Group. *Structured Assurance Case Metamodel (SACM)*, version 2.3, OMG formal/23-05-08, October 2023. [OMG SACM 2.3](https://www.omg.org/spec/SACM/2.3/).
- [7] National Aeronautics and Space Administration. *Software Assurance and Software Safety Standard*, NASA-STD-8739.8B, Section 4.4.1.2, 2022. [NASA-STD-8739.8B](https://ntrs.nasa.gov/api/citations/202200010001/main.pdf).
- [8] National Institute of Standards and Technology. *Secure Hash Standard*, FIPS PUB 180-4, 2015. [NIST FIPS 180-4](https://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.180-4.pdf).